

Call Analysis - features by release

Release	ID	Feature	Description
R0 Foundation (internal)	F1	Analytics data model	One interaction-analysis record per call; one topic record per distinct subject discussed on the call (one-to-many). Each topic carries its category, subcategory and a short description of what was discussed, used later for category discovery. The structure allows drill-down attributes to be added per topic without breaking changes. Findings attach to the interaction; group, caller type and channel are resolved through the interaction tables, not copied.
	F2	Taxonomy storage with versions and audit trail	Tables for the category / subcategory tree, client-specific sets, set versions and applicability (payer, plan, group, subgroup). Categories are immutable; a change creates a new row and a new version. Every change records who made it and when.
	F3	Taxonomy v1 agreed with the call center	First member and provider category / subcategory trees, drafted by the team and reviewed with the call center leads. Loaded as version 1 (scripted).
	F4	Client-specific sets and opt-out	CareSource Ohio PCIR and PCDR sets loaded and bound to the payer; plans that do not approve AI processing are listed for exclusion.
R1 Categorize and report	F5	Topic decomposition and categorization of every call	After summarization, each transcript is split into the distinct topics discussed (e.g. crown coverage, then find a dentist; or the same question about three different patients). Each topic is assigned a category and subcategory from the active taxonomy, plus any client-specific set that applies, and stored with the set version and a short description. Caller type comes from the interaction, not the model. Category and subcategory only; no deeper extraction yet.
	F6	Primary category per call	Each call also records the primary category, so headline reports can count calls as well as topics.
	F7	Taxonomy maintenance	A basic internal screen to view versions, add or retire categories and publish a new version. If not built in R1, taxonomy changes are scripted until it exists.
	F8	Data available in Snowflake	The new tables flow to Snowflake through the existing RCST stream.
	F9	Canned reports	Call volume and percentage by category and subcategory, selectable by group, payer or caller type. Includes CareSource Ohio PCIR and PCDR, replacing the manual monthly process.
R2 Drill down	F10	Drill-down extraction per subcategory	For each subcategory, extract the specifics that make it actionable: e.g. covered procedure - which procedure code; find a dentist - zip code search, list emailed vs read by phone; toolkit - password vs MFA. The attribute list is configured per subcategory, reviewed with the business on sample output, and stored without schema changes.
	F11	Targeted reprocessing	When an extraction definition or a taxonomy branch changes, re-run only the topics under that node or within a date range.
	F12	Snowflake tables for reporting	Denormalized tables joining topics to caller type, channel, group, subgroup and program, plus the aggregates the dashboard needs.
	F13	Interactive dashboard (pre-built)	Top categories and subcategories by volume for the last 30, 60 or 90 days; filter by caller type, group, subgroup and program (commercial vs government); drill into the extracted specifics of a subcategory. No period-over-period comparison yet.
	F14	Self-service BI	The same Snowflake tables are available to Power BI / Tableau users to build their own views.
R3 Discover and curate	F15	Category discovery	Periodic clustering over the stored topic descriptions surfaces categories and subcategories not in the tree (including what currently lands in 'other').
	F16	Proposed version and approval	A discovery run produces a proposed new taxonomy version. The business reviews, approves or rejects it; an approved version becomes the active set for new calls. Business users can also manage sets, versions and applicability directly, without a feature-team change.
	F17	Trend analysis	Month-over-month and period comparison on the dashboard, including within a single group.
	F18	Chatbot sessions	Chatbot transcripts run through the same pipeline, with a record of whether the bot fulfilled the request.
R4 Explore	F19	Chat with the data	Conversational questions over the Snowflake data via its semantic layer.
	F20	Ad-hoc mining and keyword search	User-defined filters, time windows and keyword search across topics and transcripts.
	F21	AI-suggested patterns	AI reviews a set of topics (e.g. all denture questions for a group) and suggests patterns or attributes worth capturing and reporting.
	F22	Routing and misroute analysis	Original queue vs actual group; cross-plan transfers; effect of IVR changes. Needs Genesys queue data.
	F23	Predictive use and e-mail channel	Forecast volume for large-group conversions; categorize e-mail inquiries.